

# INTRODUCTION TO BIG DATA

ECAP456

Dr. Rajni Bhalla  
Associate Professor

# Learning Outcomes



After this lecture, you will be able to

- explore concepts of distribution models
- learn types of distribution data

# Distribution Models




Run databases on a  
large cluster

# Distribution Models




Run databases on a  
large cluster



As data volumes  
increase

# Distribution Models



Bigger server to run the database

# Distribution Models

A more appealing option is


## Scale Up



## Scale Out



# Distribution Models



Aggregate orientation fits well with scaling out

# Distribution Models

## Important Benefits of Distribution Model



Handle larger quantities of data

# Distribution Models

## IMPORTANT BENEFITS OF DISTRIBUTION MODEL



Handle larger quantities of data

Process a greater read or write traffic

# Distribution Models



Network slowdowns or breakages

# Distribution Models



These are often important benefits, but they come at a cost. Running over a cluster introduces complexity - so it's not something to do unless the benefits are convincing.

# Two styles of distributing data

There are two styles of distributing data:


Sharding:

Replication

Master –slave


Peer-to-peer

# Two styles of distributing data



REPLICATION

# Two styles of distributing data



Sharding

# Single Server

# Single Server


- No distribution at all.
- Run the database on a single.
- It eliminates all the complexities that the other options introduce
- Easy for operations people to manage
- Easy for application developers to reason about.

# Single Server

**Ques: Which database works best in single-server configuration ?**

# Single Server

Ques: Which database works best in single-server configuration ?




Graph Database

# Single Server


**What is the most appropriate solution for  
processing aggregates in a single-server  
document?**

# Single Server



Key-value store

# Sharding



# Sharding

Original Table

CUSTOMER ID	FIRST NAME	LAST NAME	CITY
1	Alice	Anderson	Austin
2	Bob	Best	Boston
3	Carrie	Conway	Chicago
4	David	Doe	Denver

Vertical Shards

CUSTOMER ID	FIRST NAME	LAST NAME
1	Alice	Anderson
2	Bob	Best
3	Carrie	Conway
4	David	Doe

CUSTOMER ID	CITY
1	Austin
2	Boston
3	Chicago
4	Denver

Horizontal Shards

CUSTOMER ID	FIRST NAME	LAST NAME	CITY
1	Alice	Anderson	Austin
2	Bob	Best	Boston

HS2

CUSTOMER ID	FIRST NAME	LAST NAME	CITY
3	Carrie	Conway	Chicago
4	David	Doe	Denver

# Why Is Sharding Used?

- **Rapid responses from that server.**
- The load is balanced out nicely between servers.
- You can store the new chunks of data, called logical shards
- Can take advantage of all the compute resources across your cluster for every query.
- Scan fewer rows when responding to a query.

# Why Is Sharding Used?

- Rapid responses from that server.
- **The load is balanced out nicely between servers.**
- You can store the new chunks of data, called logical shards
- Can take advantage of all the compute resources across your cluster for every query.
- Scan fewer rows when responding to a query.

# Why Is Sharding Used?

- Rapid responses from that server.
- The load is balanced out nicely between servers.
- You can store the new chunks of data, called logical shards
- Can take advantage of all the compute resources across your cluster for every query.
- Scan fewer rows when responding to a query.

# Why Is Sharding Used?

- Rapid responses from that server.
- The load is balanced out nicely between servers.
- You can store the new chunks of data, called logical shards
- Can take advantage of all the compute resources across your cluster for every query.
- Scan fewer rows when responding to a query.

# Why Is Sharding Used?


- Rapid responses from that server.
- The load is balanced out nicely between servers.
- You can store the new chunks of data, called logical shards
- Can take advantage of all the compute resources across your cluster for every query.
- Scan fewer rows when responding to a query.

# Sharding

**How to clump the data up so that one user  
mostly gets her data from a single server ?**


# Sharding

- How to clump the data up so that one user mostly gets her data from a single server




# Sharding

- Combine data that's commonly accessed together.
- There are several factors that can help improve performance.




# Sharding

- Combine data that's commonly accessed together.
- There are several factors that can help improve performance.




# Sharding

- Place the data close to where it's being accessed.
- If you have orders for someone who lives in Boston, you can place that data in your eastern US data center




# Sharding

- Place the data close to where it's being accessed.
- If you have orders for someone who lives in Boston, you can place that data in your eastern US data center




# Sharding

- To keep the load even
- Puts Aggregate together
- Keeping its rows in lexicographic order.
- sorting web addresses based on reversed domain names




# Sharding

- To keep the load even
- Puts Aggregate together
- Keeping its rows in lexicographic order.
- sorting web addresses based on reversed domain names




# Sharding

- To keep the load even
- Puts Aggregate together
- Keeping its rows in lexicographic order.
- sorting web addresses based on reversed domain names



# Sharding

- To keep the load even
- Puts Aggregate together
- Keeping its rows in lexicographic order.
- Sorting web addresses based on reversed domain names




# Sharding

- Historically most people have done sharding as part of application logic

A to D

E to G

# Sharding




Rebalancing the Sharding

# Sharding

## Rebalancing the Sharding

- Changing the application code.
- Migrating the data.

# Auto-Sharding



Auto-sharding  
Shards mongo mongo  
mongo

# When Horizontal Sharding is Effective?

User Id	Name	Email
1	User 1	u1@nlogin.in
2	User 2	u2@nlogin.in
3	User 3	u3@nlogin.in
4	User 4	u4@nlogin.in

Original Table

User Id	Name	Email
1	User 1	u1@nlogin.in
2	User 2	u2@nlogin.in

Shard 1

User Id	Name	Email
3	User 3	u3@nlogin.in
4	User 4	u4@nlogin.in

Shard 2

Horizontal Partitioning  
or  
Database Sharding

# When Vertical Sharding is Effective?

User Id	Name	Email
1	User 1	u1@nlogin.in
2	User 2	u2@nlogin.in
3	User 3	u3@nlogin.in
4	User 4	u4@nlogin.in

Original Table

User Id	Name
1	User 1
2	User 2
3	User 3
4	User 4

Partation 1

User Id	Email
1	u1@nlogin.in
2	u2@nlogin.in
3	u3@nlogin.in
4	u4@nlogin.in

Partation 2

Vertical Partitioning


# Sharding

- Sharding is particularly valuable for performance.
- Using replication, particularly with caching, can greatly improve read performance.
- Sharding provides a way to horizontally scale writes.

# Sharding

- Improve resilience when used alone.
- Although the data is on different nodes, a node failure makes that shard's data unavailable just as surely as it does for a single-server solution.

# Master-Slave Replication



Data is replicated from master to slaves. The master services all writes; reads may come from either master or slaves.

# Master-Slave Replication

- Master-slave replication is most helpful for scaling when you have a read-intensive dataset.
- You can scale horizontally to handle more read requests by adding more slave nodes and ensuring that all read requests are routed to the slaves.


# Master-Slave Replication

- You are still, however, limited by the ability of the master to process updates and its ability to pass those updates on.
- Consequently it isn't such a good scheme for datasets with heavy write traffic, although offloading the read traffic will help a bit with handling the write load.

# Master-Slave Replication

- A second advantage of master-slave replication is **read resilience**:
- The failure of the master does eliminate the ability to handle writes until either the master is restored or a new master is appointed.

# Peer-to-Peer Replication




# Combining Sharding and Replication

master for two shards



slave for two shards




master for one shard



master for one shard  
and slave for a shard



slave for two shards






slave for one shard



Using master-slave replication  
together with sharding

# Combining Sharding and Replication



Using peer-to-peer replication  
together with sharding

# What Is the Difference between Sharding and Partitioning?

- breaking up a large data set into smaller subsets.
- sharding implies the data is spread across multiple computers while partitioning does not.
- Partitioning is about grouping subsets of data within a single database instance.
- In many cases, the terms sharding and partitioning are even used synonymously.

# What Is the Difference between Sharding and Partitioning?

- breaking up a large data set into smaller subsets.
- sharding implies the data is spread across multiple computers while partitioning does not.
- Partitioning is about grouping subsets of data within a single database instance.
- In many cases, the terms sharding and partitioning are even used synonymously.

# What Is the Difference between Sharding and Partitioning?

- breaking up a large data set into smaller subsets.
- sharding implies the data is spread across multiple computers while partitioning does not.
- **Partitioning is about grouping subsets of data within a single database instance.**
- In many cases, the terms sharding and partitioning are even used synonymously.

# What Is the Difference between Sharding and Partitioning?

- breaking up a large data set into smaller subsets.
- sharding implies the data is spread across multiple computers while partitioning does not.
- Partitioning is about grouping subsets of data within a single database instance.
- In many cases, the terms sharding and partitioning are even used synonymously.



**That's all for now...**