

# INTRODUCTION TO BIG DATA

ECAP456

Dr. Rajni Bhalla  
Associate Professor

# Learning Outcomes



After this lecture, you will be able to

- learn MapReduce types
- learn input formats and output formats in hadoop and mapreduce

# Map Reduce Types

The map and reduce functions in Hadoop MapReduce have the following general form:

map:  $(K_1, V_1) \rightarrow \text{list}(K_2, V_2)$

reduce:  $(K_2, \text{list}(V_2)) \rightarrow \text{list}(K_3, V_3)$

# Map Reduce Types

```
public void map(LongWritable key, Text value, Context context)
{ ....
.....
context.write(new Text(year), new IntWritable(airTemperature));
}

public void reduce(Text key, Iterable values, Context context)
{
.....
context.write(key, new IntWritable(maxValue));
```

# Map Reduce Types

If a combine function is used, then it is the same form as the reduce function (and is an implementation of Reducer), except its output types are the intermediate key and value types (K2 and V2), so they can feed the reduce function:

map:  $(K1, V1) \rightarrow \text{list}(K2, V2)$

combine:  $(K2, \text{list}(V2)) \rightarrow \text{list}(K2, V2)$

reduce:  $(K2, \text{list}(V2)) \rightarrow \text{list}(K3, V3)$

Often the combine and reduce functions are the same, in which case, K3 is the same as K2, and V3 is the same as V2.

# Map Reduce Types

**How input types are set?**

# Map Reduce Types

```
job.setOutputKeyClass(Text.class);  
job.setOutputValueClass(IntWritable.class);
```

# Map Reduce Types

So if K2 and K3 are the same, you don't need to call `setMapOutputKeyClass()`, since it falls back to the type set by calling `setOutputKeyClass()`.

Similarly, if V2 and V3 are the same, you only need to use `setOutputValueClass()`.

# Configuration of MapReduce in the new API

Table 7-1. Configuration of MapReduce types in the new API

Property	Job setter method	Input types		Intermediate types		Output types	
		K1	V1	K2	V2	K3	V3
Properties for configuring types:							
mapreduce.job.inputformat.class	setInputFormatClass()	.	.				
mapreduce.map.output.key.class	setMapOutputKeyClass()			.			
mapreduce.map.output.value.class	setMapOutputValueClass()			.			
mapreduce.job.output.key.class	setOutputKeyClass()				.		
mapreduce.job.output.value.class	setOutputValueClass()				.		
Properties that must be consistent with the types:							
mapreduce.job.map.class	setMapperClass()	.	.	.	.	.	.
mapreduce.job.combine.class	setCombinerClass()			.	.	.	.
mapreduce.job.partitioner.class	setPartitionerClass()			.	.	.	.
mapreduce.job.output.key.comparator.class	setSortComparatorClass()			.			
mapreduce.job.output.group.comparator.class	setGroupingComparatorClass()			.			
mapreduce.job.reduce.class	setReducerClass()			.	.	.	.
mapreduce.job.outputformat.class	setOutputFormatClass()			.	.	.	.

# Configuration of MapReduce types in the old API

Table 7-2. Configuration of MapReduce types in the old API

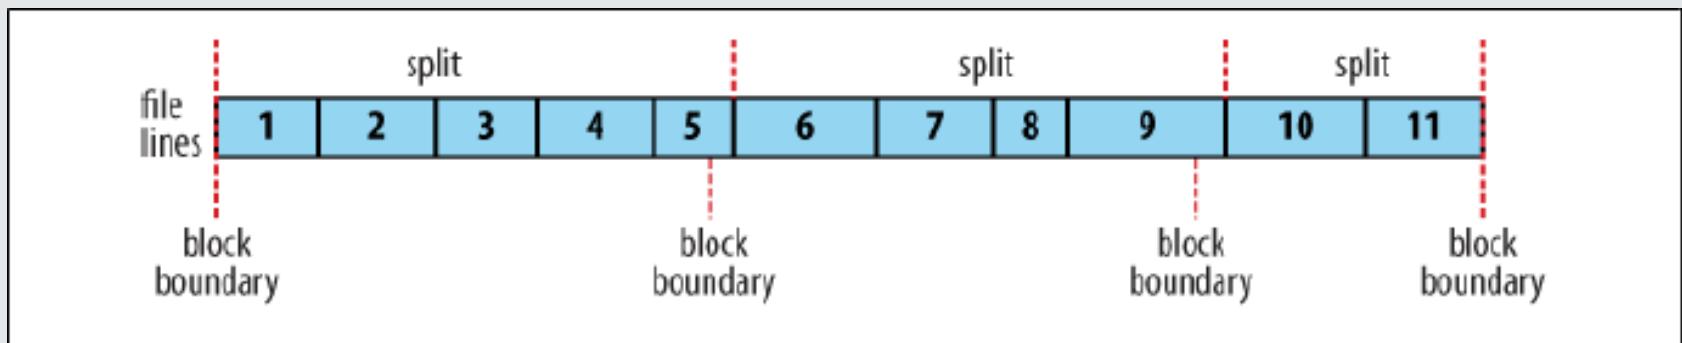
Property	JobConf setter method	Input types		Intermediate types		Output types	
		K1	V1	K2	V2	K3	V3
Properties for configuring types:							
mapred.input.format.class	setInputFormat()			.	.		
mapred.mapoutput.key.class	setMapOutputKeyClass()				.		
mapred.mapoutput.value.class	setMapOutputValueClass()				.		
mapred.output.key.class	setOutputKeyClass()				.		
mapred.output.value.class	setOutputValueClass()				.		
Properties that must be consistent with the types:							
mapred.mapper.class	setMapperClass()	.	.	.	.	.	.
mapred.map.runner.class	setMapRunnerClass()	.	.	.	.	.	.
mapred.combiner.class	setCombinerClass()			.	.		
mapred.partitioner.class	setPartitionerClass()			.	.		
mapred.output.key.comparator.class	setOutputKeyComparatorClass()			.			
mapred.output.value.groupfn.class	setOutputValueGroupingComparator()			.			
mapred.reducer.class	setReducerClass()			.	.	.	.
mapred.output.format.class	setOutputFormat()			.	.	.	.

# Input Formats

Hadoop can process many different types of data formats, from flat text files to databases

# Input Formats

## The Relationship Between Input Splits and HDFS Block



# Text Input

So a file containing the following text:

```
On the top of the Crumpetty Tree  
The     Quangle     Wangle     sat,  
But his face you could not see,  
On account of his Beaver Hat.
```

# Text Input

So a file containing the following text:	The records are interpreted as the following key-value pairs:
On the top of the Crumpetty Tree	(0, On the top of the Crumpetty Tree)
The Quangle Wangle sat,	(33, The Quangle Wangle sat,)
But his face you could not see,	(57, But his face you could not see,)
On account of his Beaver Hat.	(89, On account of his Beaver Hat.)

# Text Input

**KeyValueTextInputFormat**

# Text Input

`mapreduce.input.keyvaluelinerecodreader.key.valu  
e.separator` property

Or

`key.value.separator.in.input.line`  
in the old API

# Text Input

Consider the following input file, where → represents a (horizontal) tab character:

line1→On the top of the Crumpetty Tree

line2→The Quangle Wangle sat,

line3→But his face you could not see,

line4→On account of his Beaver Hat.

# Text Input

**NLineInputFormat**

# Text Input

On the top of the Crumpetty Tree

The Quangle Wangle sat,

But his face you could not see,

On account of his Beaver Hat.

# Text Input

If, for example, N is **two**, then each split contains two lines.

One mapper will receive the first two key-value pairs:

(0, On the top of the Crumpetty Tree)

(33, The Quangle Wangle sat,)

# Text Input

And another mapper will receive the second two key-value pairs:

(57, But his face you could not see,)

(89, On account of his Beaver Hat.)

# Text Input

## XML

set your input format to StreamInputFormat and  
set the stream.recordreader.class property to  
org.apache.hadoop.streaming.StreamXmlRecordReader  
to use xml as an input format.

# Text Input

Binary Input :

`SequenceFileInputFormat`

# Text Input

Binary Input :

`SequenceFileInputFormat`

`SequenceFileAsTextInputFormat`

# Text Input

Binary Input :

SequenceFileInputFormat

SequenceFileAsTextInputFormat

SequenceFileAsBinaryInputFormat

# Text Input

## SequenceFileInputFormat

Hadoop's sequence file format stores sequences of binary key-value pairs. Sequence files are well suited as a format for MapReduce data since they are splittable they support compression as a part of the format.

# Text Input

## SequenceFileAsTextInputFormat

SequenceFileAsTextInputFormat is a variant of SequenceFileInputFormat that converts the sequence file's keys and values to Text objects.

# Text Input

## **SequenceFileAsBinaryInputFormat**

`SequenceFileAsBinaryInputFormat` is a variant of `SequenceFileInputFormat` that retrieves the sequence file's keys and values as opaque binary objects.

# Text Input

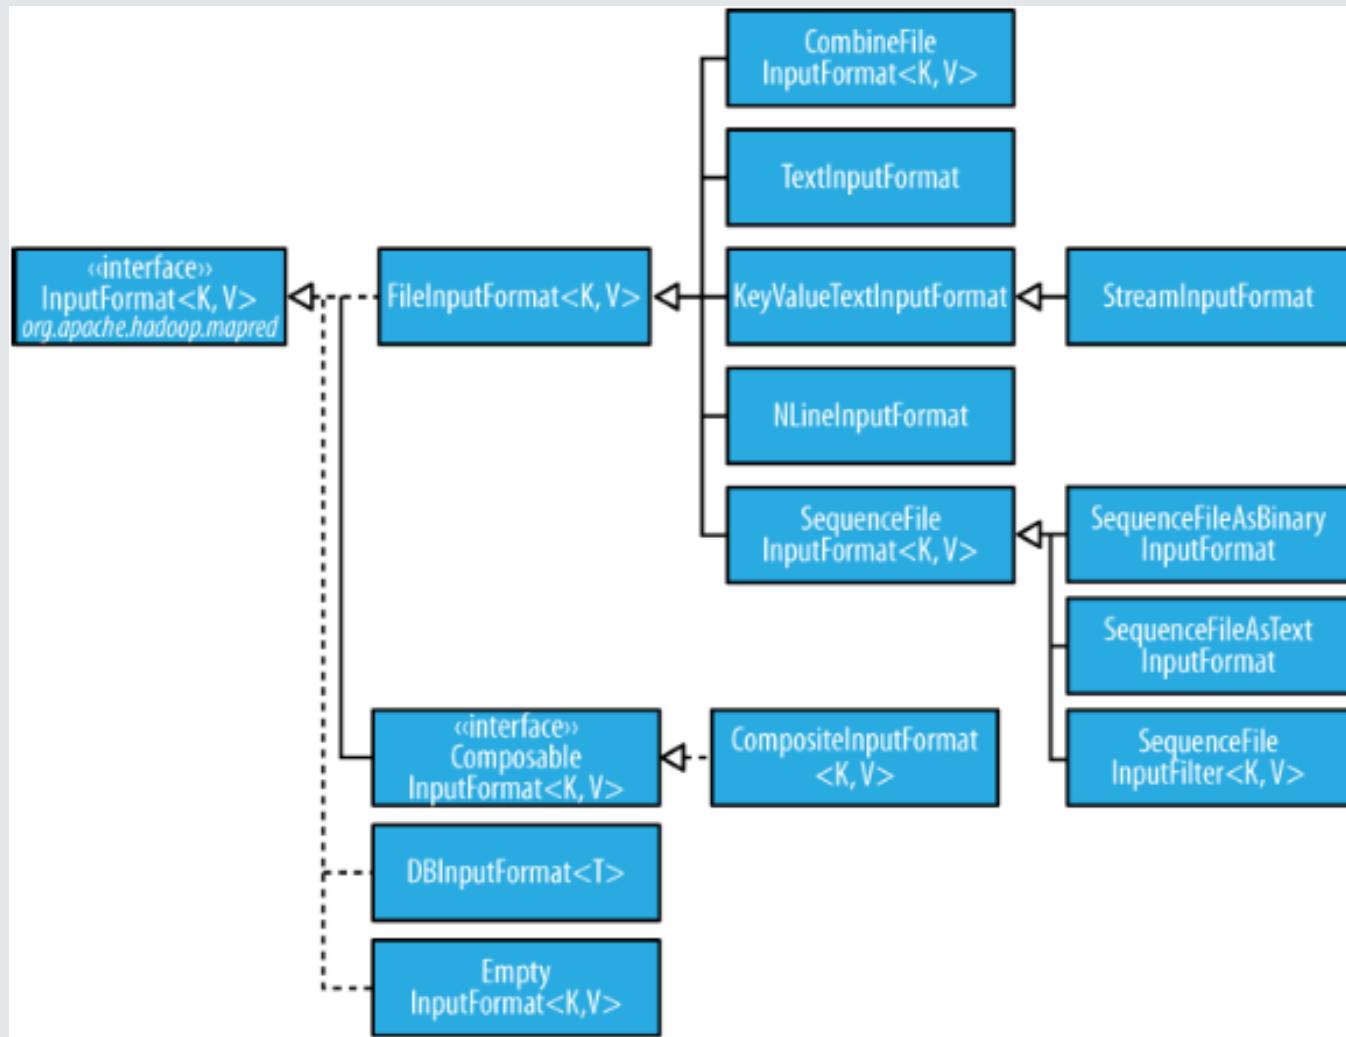
## Multiple Inputs

The input to a MapReduce job may consist of multiple input files. This case is handled elegantly by using the `MultipleInputs` class.

# Text Input

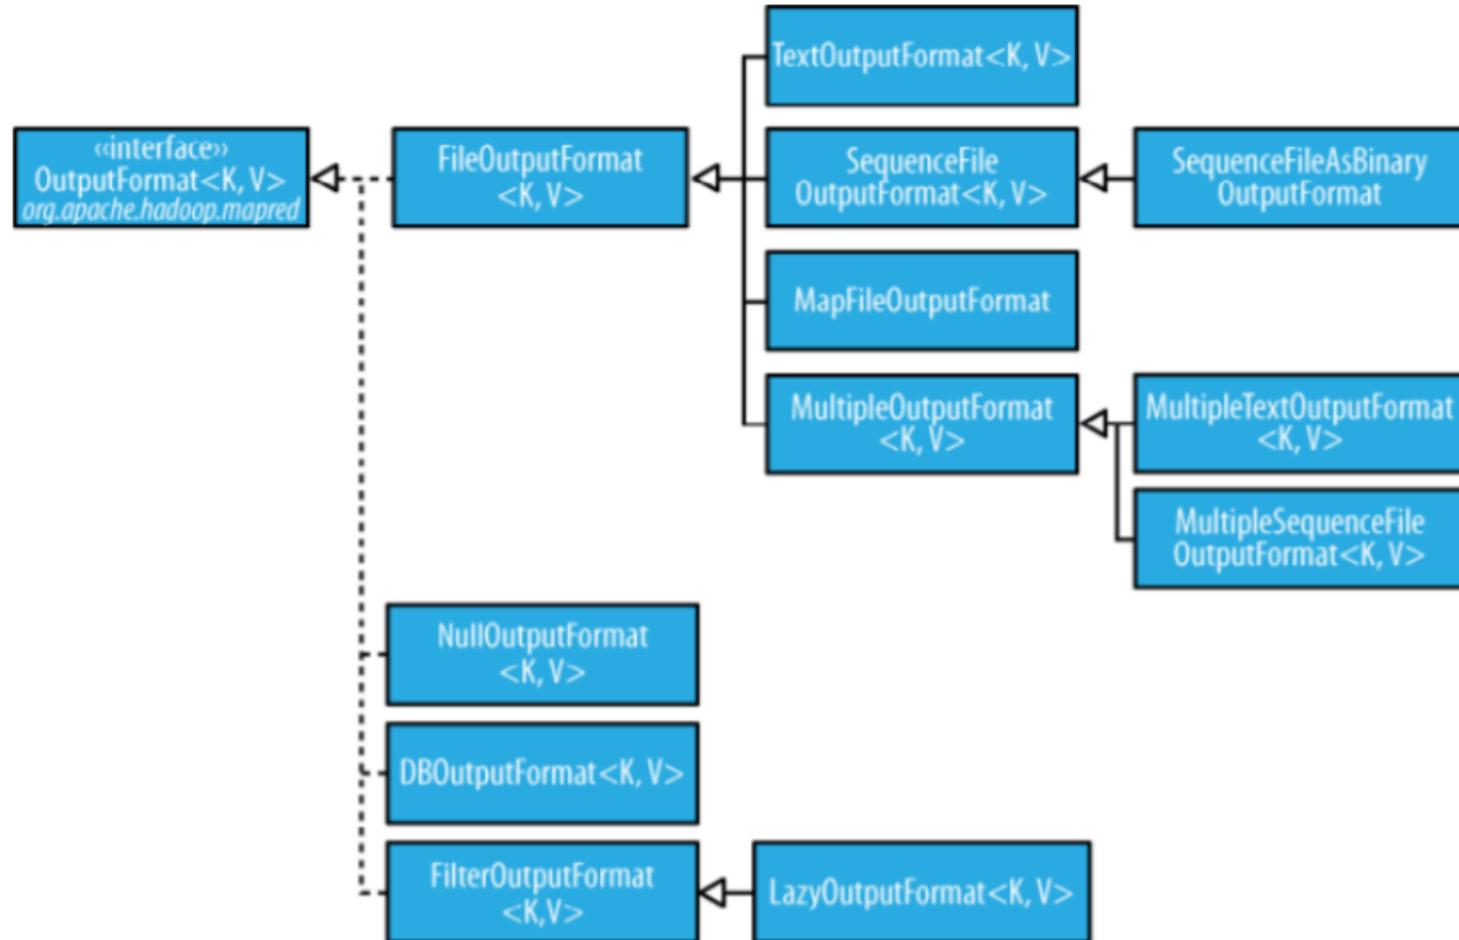
```
MultipleInputs.addInputPath(job, ncdclnputPath,  
TextInputFormat.class,  
MaxTemperatureMapper.class);  
  
MultipleInputs.addInputPath(job, metOfficeInputPath,  
TextInputFormat.class,  
MetOfficeMaxTemperatureMapper.class);
```

# Database Input



Input format class hierarchy

# Database Input



Output format class hierarchy

# Text Output

- The default output format, `TextOutputFormat`, writes records as lines of text.
- Its keys and values may be of any type, since `TextOutputFormat` turns them to strings by calling `toString()` on them.
- Each key-value pair is separated by a tab character. The counterpart to `TextOutputFormat` for reading in this case is `KeyValueTextInputFormat`.

# Binary Output

- SequenceFileOutputFormat

# Binary Output

- SequenceFileOutputFormat
- SequenceFileAsBinaryOutputFormat

# Binary Output

- SequenceFileOutputFormat
- SequenceFileAsBinaryOutputFormat
- MapFileOutputFormat

# Hadoop InputFormat

## Hadoop InputFormat

- Describes the input-specification for execution.
- Describes how to split up and read input files.
- InputFormat is the first step.
- Responsible for creating the input splits and dividing them into records.
- Input files reside in HDFS.

# Hadoop InputFormat

Hence, In MapReduce, InputFormat class is one of the fundamental classes which provides below functionality:

- It selects the files or other objects for input.
- It also defines the Data splits.
- It defines the RecordReader.

# Hadoop InputFormat

## How we get the data from Mapper?

```
public abstract class InputFormat<K, V>
{
    public abstract List<InputSplit> getSplits(JobContext context)
        throws IOException, InterruptedException;
    public abstract RecordReader<K, V>
        createRecordReader(InputSplit split,
                           TaskAttemptContext context) throws IOException,
        InterruptedException;
}
```

# Hadoop InputFormat

Types of InputFormat in MapReduce

FileInputFormat

TextInputFormat

Key

Value

KeyValueTextInputFormat

SequenceFileInputFormat

SequenceFileAsTextInputFormat

SequenceFileAsBinaryInputFormat

NlineInputFormat

DBInputFormat

# Hadoop InputFormat

Types of  
InputFormat  
in MapReduce

## 1. FileInputFormat

# Hadoop InputFormat

Types of  
InputFormat  
in MapReduce

1. FileInputFormat
2. TextInputFormat

# Hadoop InputFormat

Types of  
InputFormat  
in MapReduce

1. `FileInputFormat`
2. `TextInputFormat`
3. `KeyValueTextInputFormat`

# Hadoop InputFormat

Types of  
InputFormat  
in MapReduce

1. `FileInputFormat`
2. `TextInputFormat`
3. `KeyValueTextInputFormat`
4. `SequenceFileInputFormat`

# Hadoop InputFormat

Types of  
InputFormat  
in MapReduce

1. `FileInputFormat`
2. `TextInputFormat`
3. `KeyValueTextInputFormat`
4. `SequenceFileInputFormat`
5. `NlineInputFormat`

# Hadoop InputFormat

Types of  
InputFormat  
in MapReduce

1. **FileInputFormat**
2. **TextInputFormat**
3. **KeyValueTextInputFormat**
4. **SequenceFileInputFormat**
5. **NlineInputFormat**
6. **DBInputFormat**

# Hadoop InputFormat

Types of  
InputFormat  
in MapReduce

1. `FileInputFormat`
2. `TextInputFormat`
3. `KeyValueTextInputFormat`
4. `SequenceFileInputFormat`
5. `NlineInputFormat`
6. `DBInputFormat`
7. `Sequence FileAs BinaryInputFormat`



**That's all for now...**